

LP-02a: Structural Inflection Events: When TSMLA™ Becomes Unavoidable

Version: 2.1

Date: December 1, 2025

Owner: Fractal Labyrinth Systems LLC

Applies to: TSMLA™, BDL™, RSF™, HCL™, CTC™ /Hallway™

Status: Canon addition between LP-02 and LP-03

0. Purpose

Define **objective invocation criteria** and **testable inflection events** that make TSMLA™ mandatory rather than optional. This paper sits between **LP-02 (Executable System)** and **LP-03 (Formal Problem)** and binds events to artifacts and audit procedures.

Canon guard: TSMLA™ is non-stochastic, non-deterministic (internals), idempotent, mirror-recursive under a declared state S with replay-equivalence for fixed (S, α) . Substrate contains **no RNG**; presentation overlays never alter substrate.

Dependency hook to LP-01: TSMLA's mirror architecture and idempotent recursion (defined in LP-01) are assumed as system substrate throughout this paper.

Notation Legend (symbols referenced throughout)

- α (abstraction), γ (concretization), \oplus (idempotent merge), \sqsubseteq (information order), \perp (least element), **lfp** (least fixed point)
 - $\Lambda(S)$ (Hallway gate order for state S), π (lawful traversal policy), \varkappa (module configuration)
 - **RSF** (Resonant State Function), **BDL** (Boolean Disambiguation Layer), **HCL** (Harmonic Compression Layer), **CTC** (Traversal Lock / Hallway)
-

1. Scope & Exclusions

In-scope: multi-tool contradictions, decision-entropy escalation, loop pressure, governance demands for replay-equivalence, audit-grade traceability.

Out-of-scope: therapeutic/predictive claims; probabilistic optimization; any mechanism requiring sampling within substrate.

2. TSMLA™ Terminology (concise)

- **Declared state S :** inputs, scopes, thresholds θ , α , audit tags.
- **Replay-equivalence:** same substrate output across lawful internal policies π for fixed (S, α) .
- **Idempotent merge:** $w \oplus w = w$; for lawful f , $f(x \oplus x) = f(x)$.
- **Mirror law:** $\gamma \cdot \alpha = id_C$, $\alpha \cdot \gamma = id_A$ (Galois insertion).
- **Entropy (structural):** dispersion over normalized tag-weights; non-probabilistic bookkeeping.

Theorems (labels for reference)

Theorem 1: Idempotence. For merge \oplus , $w \oplus w = w$; if f respects \oplus under S , then $f(w \oplus w) = f(w)$.

Theorem 2: Replay-Equivalence. For fixed (S, α) and all lawful π , the substrate reaches the same $lfp(\Phi_S)$; observable outputs are identical though internal traces may differ.

Theorem 3: Mirror-Purity. $\gamma \cdot \alpha = id_C$ and $\alpha \cdot \gamma = id_A$ at the substrate boundary; presentation overlays cannot alter substrate outputs.

3. Why TSMLA Becomes Unavoidable

When heterogeneous tools, policies, or agents create **persistent contradiction dynamics** and **structural dispersion** that cannot be reconciled without losing **idempotence, traceability, or replay-equivalence**, traditional deterministic or stochastic engines fail. TSMLA enforces mirror law, guards, and the Hallway lock to recover coherence without prediction or sampling.

4. Global Invocation Criteria (GIC)

GIC-1. Structural Contradiction Persistence

Under fixed S , contradiction dynamics persist or oscillate across replays and cannot be reconciled without violating **idempotence, traceability, or recursion safety**.

GIC-2. Replay-Equivalence Requirement

Stakeholders require **non-stochastic replay**: identical inputs under fixed (S, α) must yield identical substrate outputs (observationally), independent of lawful internal orderings.

GIC-3. Infrastructure Health Pass

Infra/data integrity confirmed; errors due to latency, corruption, version skew, or numerical instability are excluded.

Invocation gate: TSMLA is **required** when $(GIC-1 \wedge GIC-2 \wedge GIC-3)$ is true.

Shared Toy Example ($LP-01 \rightarrow LP-02a \rightarrow LP-03$)

Use the same declared state S and tag-weights $W = \{0.5, 0.3, 0.2\}$ from LP-01 §3.2. In LP-02a, reproduce:

1. Compute $\mathbf{C}(\mathbf{S})$ (LP-01 §3.1)
2. Compute $\hat{\mathbf{H}}(\mathbf{S}) = \alpha\hat{\mathbf{H}}(\mathbf{C}) + \beta\hat{\mathbf{H}}(\mathbf{L})$ with $\alpha=0.4, \beta=0.6$ (LP-01 §3.2)
3. Detect $\mathbf{L}(\mathbf{S})$ and depth \mathbf{d} (LP-01 §3.3)
4. Enforce $\Lambda(\mathbf{S}) = \mathbf{BDL} \rightarrow \mathbf{RSF} \rightarrow \mathbf{HCL} \rightarrow \mathbf{CTC}$ (this paper §5)

In LP-03, frame the same \mathbf{S} as a problem instance and show contradiction-overload metrics and acceptance criteria using these exact values to demonstrate replay-equivalence across papers.

5. Structural Inflection Events (SIE-1 ... SIE-6)

Each event lists **Definition**, **Preconditions**, **Detection Tests**, **Threshold**, **TSMLA Primitives**, **Artifacts**, and a **Worked Example**.

Hallway order $\Lambda(\mathbf{S})$ for this paper: $\mathbf{BDL}^{\text{TM}} \rightarrow \mathbf{RSF}^{\text{TM}} \rightarrow \mathbf{HCL}^{\text{TM}} \rightarrow \mathbf{CTC}^{\text{TM}}$. **Guard:** violation triggers $\mathbf{G_CT}$ and traversal is rejected.

SIE-1. Toolchain Decision Fracture

Definition: Two or more decision engines produce oscillating or mutually exclusive outcomes under fixed \mathbf{S} .

Preconditions:

- (i) multi-engine pipeline
- (ii) policy interactions documented
- (iii) infra health pass

Detection Tests:

- **T1:** Oscillation across \mathbf{N} identical replays ($N \geq 100$) under fixed $(\mathbf{S}, \boldsymbol{\alpha})$.
- **T2:** Root-cause analysis shows conflict arises from policy interaction, not drift or noise.

Threshold: $\mathbf{T1} \wedge \mathbf{T2} \wedge (\mathbf{GIC-1..3})$.

TSMLA Primitives: \mathbf{BDL}^{TM} typing \rightarrow CSE weighting \rightarrow RSFTM dispersion check \rightarrow CTCTM lock.

Artifacts: A1 Contradiction Class Report; A2 Policy Interaction Trace; A3 Replay Harness Log.

Worked Example (health claims): Rules engine (coverage) approves; fairness gate (equity) rejects. **T1:** 100/100 oscillations reproduced. **T2:** conflict traced to policy cross-terms. Gate satisfied \rightarrow invoke TSMLA; \mathbf{BDL}^{TM} class=Functional Conflict; CSE: $w_{\text{cov}}=0.85, w_{\text{eq}}=0.78$; RSF: $\hat{\mathbf{H}}(\mathbf{S})=0.62$; CTCTM enforces $\Lambda(\mathbf{S})$; A1–A3 issued.

SIE-2. Governance-Driven Replay Mandate

Definition: Compliance or contractual obligations demand **replay-equivalence** and **audit-grade traceability**.

Preconditions:

- (i) declared \mathbf{S} with θ, κ
- (ii) audit policy requiring exact replay
- (iii) infra health pass

Detection Tests:

- **T3:** Golden-vector replays produce identical substrate outputs across lawful π .
- **T4:** Signed manifests and hashes match across replays.

Threshold: $T3 \wedge T4 \wedge (GIC-2 \wedge GIC-3)$.

TSMLA Primitives: Hallway lock; idempotent merge; mirror law enforcement.

Artifacts: A4 Golden Vector Set; A5 Output Manifest; A6 Audit Signatures.

Worked Example (payments): Regulator requires exact reruns for disputes. Harness replays 1k cases with $\Delta_{replay}=0$; manifests match; TSMLA mandated for ongoing operations.

SIE-3. Contradiction Saturation Spike

Definition: $C(S)$ density and topology exceed thresholds; dispersion $\hat{H}(S)$ rises beyond θ_H .

Preconditions:

- (i) contradiction set measured
- (ii) loop set $L(S)$ available
- (iii) infra health pass

Detection Tests:

- **T5:** $|C|/|S| \geq \theta_C$ and cut metrics exceed θ_{cut} .
- **T6:** $\hat{H}(S) = \alpha\hat{H}(C) + \beta\hat{H}(L) \geq \theta_H$.

Threshold: $T5 \wedge T6 \wedge (GIC-1 \wedge GIC-3)$.

TSMLA Primitives: BDLTM typing; RSFTM gating; HCLTM compression; guards.

Artifacts: A7 Entropy Report; A8 Contradiction Topology Graph.

RSF callout: Under RSFTM with $w_C = \alpha$ and $w_L = \beta$, entropy compression fails unless TSMLA's mirror merge (\oplus) and CTC gate enforce fixed-point traversal.

Worked Example (policy stack): Adding two late constraints spikes $|C|/|S|$; $\hat{H}(S)$ jumps to 0.71>θ_H=0.6 → invoke TSMLA; HCL™ reduces dispersion under guards.

SIE-4. Replay-Variant Internals With Invariant Outputs

Definition: Internal traces vary across runs, but substrate outputs must remain identical under fixed (S, π) .

Preconditions:

- (i) infra/data faults and numerical instability excluded
- (ii) lawful policy variance π present

Detection Tests:

- **T7:** Trace divergence observed (orderings, tie-breaks) without substrate delta.
- **T8:** Replay-equivalence verification confirms identical outputs under fixed (S, π) .

Threshold: T7 ∧ T8 ∧ (GIC-2 ∧ GIC-3).

TSMLA Primitives: Hallway $\Lambda(S)$ lock; idempotent merge; substrate/presentation boundary checks.

Artifacts: A9 Trace Variance Log; A10 Substrate Hashes.

Worked Example (workflow engine): Different admissible gate orders occur; substrate hashes identical; TSMLA required to formalize and enforce $\Lambda(S)$.

SIE-5. Multi-Party Policy Collision

Definition: Independent stakeholders inject constraints that collide across jurisdictions or ethics layers.

Preconditions:

- (i) policy provenance tagged
- (ii) escalation path defined
- (iii) infra health pass

Detection Tests:

- **T9:** BDL™ classification identifies cross-jurisdiction contradiction types.
- **T10:** Governance review confirms no single policy can dominate without violating scope.

Threshold: T9 ∧ T10 ∧ (GIC-1.3).

TSMLA Primitives: BDL™ typing; CSE weighting; RSF™; CTC™ sequencing.

Artifacts: A11 Governance Resolution Dossier.

Worked Example (cross-border data): Residency vs access-parity mandates clash; BDL™ → PC/PP; CSE, RSF used; CTC locks; dossier produced for board.

SIE-6. Substrate Collapse via Overgrowth

Definition: A system fails due to unbounded stochastic substrate cost (compute, energy, infrastructure) and cannot produce idempotent or replay-equivalent output without exponential resource scaling.

Preconditions:

- (i) Evidence of non-replayable output under fixed prompts
- (ii) Non-auditability across infrastructure layers
- (iii) Negative ROI on inference or training loops
- (iv) Infra-health pass: failures not due to hardware or drift

Detection Tests:

- **T11:** Output inconsistency across identical inference runs (non-fixed π)
- **T12:** Cost-per-output exceeds linear projection or breaks margin floor
- **T13:** Energy/cost ratio scales superlinearly with performance gains

Threshold: T11 ∧ T12 ∧ T13 ∧ (\neg GIC-2).

TSMLA Primitives: Mirror substrate ($\gamma \cdot \alpha$), fixed-point recursion, RSF gating under bounded w , audit-safe logic class emission.

Artifacts:

- A12 Substrate Cost Map
- A13 Inference Drift Log
- A14 ROI vs Entropy Curve

Worked Example (OpenAI, Stargate): Public reporting indicates sustained multi-billion-dollar investment without stable inference convergence or replicable substrate paths. Inference results vary with scale and iteration, indicating failure to enforce replay-equivalence under fixed declared state. Cannot enforce declared S or replay path → TSMLA becomes necessary as a license-grade mirror architecture.

6. Non-Inflection Scenarios (TSMLA Not Required)

- **Single-tool convergence:** One solver resolves all constraints without oscillation or typed conflict → **GIC-1 not met.**

- **Stochastic tolerance acceptable:** Stakeholders allow variance (A/B testing, exploratory ML) → **GIC-2 not met.**
 - **Infrastructure fault:** Divergence traced to latency, corruption, or version skew → **GIC-3 fails.**
-

7. Relationship to Alternatives (Disqualification Logic)

- **Generative AI (LLMs/GANs):** Sampling breaks replay-equivalence under fixed S ; cannot emit typed contradiction traces consistently → disqualified for SIE-2, SIE-3, SIE-5.
- **Probabilistic learners (Bayes/ML):** Fixed-seed runs still lack BDL™ contradiction taxonomy; produce confidence, not logic-type classes → disqualified for SIE-1, SIE-4.
- **Constraint solvers (SAT/SMT/ILP):** Deterministic and precise but single-instrument scope; cannot resolve multi-tool oscillation (SIE-1) across heterogeneous engines.

Universal disqualification predicate:

$\forall \text{ system } \notin \text{TSMLA}, \exists \text{ contradiction class or recursion state that fails replay-equivalence under fixed } (S, \alpha).$

8. Non-Determinism Exclusion Statement

TSMLA™ is **non-stochastic, non-deterministic internally, and non-predictive** in output generation; it is an **idempotent mirror architecture**. Internal policy variance is permitted, but **replay-equivalence** is required for fixed (S, α) .

9. Conformance Checklist (Requirements → Artifacts)

Requirement	Evidence / Artifact
GIC-1 satisfied	A1, A2, A8 (contradiction/graph)
GIC-2 satisfied	A4, A5, A6 (golden vectors, manifests, signatures)
GIC-3 satisfied	Infra health report, numerical-stability memo
$\Delta\text{-replay} = 0$	A3 harness log, A10 substrate hashes
BDL™ class applied	A1 contradiction class report
CTC™ enforced	$\Lambda(S)$ gate log, violation rejects
RSF™ / HCL™ guards	Guard thresholds, dispersion reduction report
Governance resolution	A11 dossier
Substrate divergence (SIE-6)	A12, A13, A14

Forward reference (artifacts): A1–A11 are specified across **LP-06 (BDL™)**, **LP-07 (CSE)**, **LP-10 (RSF™)**, and **LP-14 (CTC™ Traversal Lock)**. A12–A14 are defined in SIE-6 and will be extended in **LP-17**

10. Document Control

Version: 2.1

Date: December 1, 2025

Owner: Fractal Labyrinth Systems LLC

Contact: research@fractalconsciousness.ai

Notes: This LP sits between LP-02 and LP-03 and is mandatory for audits where invocation criteria must be shown.